

Supplementary Material

Architecture and Preprocessing Choices under Dataset Shift in Single-Channel EEG Sleep Staging: A Controlled Evaluation on Sleep-EDF and SHHS1

Ngo Nhat Quan, Pham Thai Son, Tri Nguyen Ho Duy and Tri Nguyen

This supplement provides model specifications, training settings and supporting statistical analyses. Seed 42 is primary, seed 123 is a post-protocol sensitivity repeat, SHHS1 E1/E2 analyses are secondary on an opened cohort, and SHHS1 E3–E2/E4 analyses are post-hoc. The E4 extension now includes both seeds 42 and 123 on that same cohort; seed-specific results are kept separate and do not alter the original primary comparison family. E6 estimates mean and standard deviation from each complete unlabelled target recording; it is label-free but transductive at recording level, not purely inductive zero-shot inference.

Supplementary Methods S1: model and training details

Architecture and feature construction

Each encoder input is a 30-s epoch with 3,000 samples. The control encoder contains 15 separately trained CNNs: three context groups (current, previous and next; C/P/N) crossed with five class-specialised loss functions. Every CNN has a five-class output, not a binary head. Concatenating its softmax outputs in C/P/N order gives 75 features per central epoch. The P and N inputs are shifted within each recording and are trained against the central epoch’s label; the first or last signal epoch is replicated at a recording edge. No context crosses recording boundaries. These are softmax-output features; their use does not imply that probability calibration has been established.

The ResNet-1D encoder instead produces a 128-dimensional, pre-dropout representation of the current epoch. Layer specifications are given in Table S1. Encoders are trained and selected before sequence training. Their selected checkpoints are loaded in evaluation mode, and features are cached without gradients. E1 reuses the E0 encoder and its features within the same fold and seed; its sequence training settings differ from E0 (Table S2). E1–E0 is therefore a sequence-model configuration comparison, not an architecture-only effect. E2–E1 replaces the feature/context package, including encoder training and selection, while retaining the specified TCN recipe.

Table S1: Layer specifications. Convolution entries give input/output channels, kernel, stride and padding; d denotes dilation. All final classifiers have five outputs.

Component	Layers and operations
One control CNN	Conv1d $1 \rightarrow 10$, $k = 55$, $s = 1$, $p = 27$, no bias; ReLU; batch normalisation; max-pool $k = s = 16$, ceiling mode. Then Conv1d $10 \rightarrow 5$, $k = 25$, $s = 1$, $p = 12$, no bias; ReLU; batch normalisation; the same max-pool. Flatten $5 \times 12 = 60$; linear $60 \rightarrow 10$; ReLU; linear $10 \rightarrow 5$.
ResNet stem	Conv1d $1 \rightarrow 32$, $k = 50$, $s = 5$, $p = 25$, no bias; batch normalisation; ReLU; max-pool $k = 3$, $s = 2$, $p = 1$.
ResNet blocks	Three residual blocks: $32 \rightarrow 64$ ($s = 1$), $64 \rightarrow 128$ ($s = 2$), $128 \rightarrow 128$ ($s = 2$). Each has two bias-free $k = 7$, $p = 3$ convolutions, with the stated stride in the first and stride 1 in the second; batch normalisation after each, ReLU after the first and after residual addition. Each shortcut uses a bias-free $k = 1$ projection with the first convolution’s stride, followed by batch normalisation.
ResNet output	Global average pooling gives 128 features. During encoder training only, dropout 0.5 precedes the linear $128 \rightarrow 5$ classifier. Cached features are taken before dropout.
BiLSTM	One bidirectional LSTM layer, input 75, hidden size 128 per direction; linear $256 \rightarrow 5$ output. Packed sequences exclude batch padding from recurrence.
TCN	Linear input projection 75 or $128 \rightarrow 128$; layer normalisation; GELU. Six residual blocks, each with two bias-free $128 \rightarrow 128$ convolutions, $k = 3$, $s = 1$, $p = d$, with $d \in \{1, 2, 4, 8, 16, 32\}$. Each convolution is followed by per-epoch layer normalisation, GELU and dropout 0.2. The shortcut is identity. Final dropout 0.2 and linear $128 \rightarrow 5$.

Training and checkpoint selection

All stages use the same subject partitions, with validation at the end of each epoch, but not the same optimisation or selection rule. For a CNN specialised to class c , training counts n_c and $N = \sum_c n_c$ determine weighted five-class cross-entropy: the weight for c is $0.5/(n_c/N)$ and each other class has weight $0.5/((N - n_c)/N)$. These training-derived weights are also used

to evaluate its validation loss. ResNet uses weights $N/(5n_c)$, again computed only from training subjects. Sequence-model cross-entropy is unweighted over valid labelled epochs. CNN branch seeds are the campaign seed plus the branch index in C/P/N-stage order; ResNet and sequence training use the campaign seed. No label smoothing, data augmentation or learning-rate scheduler is used.

Table S2: Specified training recipes. Patience counts end-of-epoch validation checks without an improvement. E1 reuses the E0 encoder; all other reported encoder/sequence pairs are fitted for their own input variant.

Setting	CNN15 (E0/E1)	BiLSTM (E0)	ResNet (E2–E6)	TCN (E1–E6)
Optimiser	Adam	Adam	AdamW	Adam
Learning rate	10^{-3}	10^{-2}	10^{-3}	5×10^{-4}
Weight decay	0	0	10^{-4}	0
Batch size	64 epochs	4 recordings	64 epochs	8 recordings
Maximum epochs	1,000	1,000	40	300
Patience	10	10	8	30
Selected checkpoint	Minimum weighted validation loss	Maximum validation macro-F1	Maximum validation macro-F1	Maximum validation macro-F1
Gradient-norm clipping	None	1.0	None	1.0

Whole retained recordings are the sequence units for training and inference. Batch padding and unknown/movement labels are excluded from the loss and metrics; unknown epochs retain their positions in the signal sequence. The TCN masks padded positions after the input projection and each convolutional sub-block. Its theoretical receptive field is $1 + 2(3 - 1)(1 + 2 + 4 + 8 + 16 + 32) = 253$ feature epochs, spanning up to 126 epochs on each side, subject to recording boundaries. This is not a causal architecture. E1 also receives the explicit neighbouring signal epochs through its CNN features. The 100-epoch input used for timing below is a benchmark shape, not a training-window length.

Signal operations and evaluation windows

Signal values are read in physical microvolts. Sleep-EDF uses Fpz–Cz at 100 Hz. SHHS1 uses C4–A1 and is resampled continuously from 125 to 100 Hz with polyphase resampling (up/down = 4/5, Kaiser window $\beta = 5$, constant padding). Table S3 specifies operations after this cohort-specific loading/resampling step. Filtering uses a fourth-order Butterworth design with 0.5–30 Hz cutoffs, implemented as second-order sections with forward–backward filtering. It operates on the continuous recording before epoch extraction and trimming.

Table S3: Input variants and fitted components. Operations are listed in execution order; all model weights are learned on Sleep-EDF only.

ID	Signal operations	Encoder and sequence configuration
E0	Raw microvolt signal	CNN15 C/P/N, then BiLSTM
E1	Raw microvolt signal	Reused E0 CNN15 C/P/N, then TCN
E2	Raw microvolt signal	Current-epoch ResNet, then TCN
E3	Band-pass; clip at $\pm 800 \mu\text{V}$; divide by 100	Current-epoch ResNet, then TCN
E4	Band-pass	Current-epoch ResNet, then TCN
E5	Band-pass; clip at $\pm 800 \mu\text{V}$	Input-identity audit; excluded from performance analysis
E6	Band-pass; clip at $\pm 800 \mu\text{V}$; recording-wise z -score	Current-epoch ResNet, then TCN

For E6, mean and population standard deviation (ddof = 0) are calculated over the entire filtered and clipped recording, including portions outside the later evaluation window. The implementation rejects zero or non-finite standard deviation rather than adding an epsilon; non-finite input/output signals are also rejected. Processing uses float64 for filtering and normalisation and stores model inputs as float32. On SHHS1 this is label-free target-record normalisation with a transductive dependency; it does not update model weights. E3 and E6 are trained separately on their respective source input variants, so their comparison is not an amplitude-only intervention on one frozen classifier.

The evaluation window extends up to 30 minutes before the first and after the last reference sleep epoch (N1, N2, N3 or REM), clipped to recording limits. Original stage 3/4 labels are merged as N3; unknown/movement epochs are unscored. This reference-dependent trimming convention and the later boundary diagnostics must be distinguished from label-free model fitting/inference. Sleep-EDF metrics use out-of-fold predictions, with each subject evaluated by a model not trained on that subject. SHHS1 uses the arithmetic mean of class probabilities from all ten source-fold models before argmax. Cross-cohort score differences therefore compare distinct cohorts and evaluation settings, not an isolated effect of dataset shift.

Runtime measurement

The archived forward-pass benchmark used the selected fold-00, seed-42 checkpoints on a Tesla V100-PCIE-16GB GPU, Python 3.10.13, PyTorch 2.5.1+cu124, CUDA 12.4 and cuDNN 9.1. It used float32 synthetic input of shape (1, 100, 1, 3000), 20 warm-up passes and 100 timed repetitions in each of three rounds, with model order varied across rounds and GPU synchronisation around timing. Both encoder and sequence forward passes are included; disk I/O, preprocessing, feature-cache serialization and training are excluded. Latency is the median of the recorded passes; memory is the maximum peak allocated GPU memory, distinct from reserved or incremental memory. These measurements characterise offline computation, not the time until future EEG needed by a bidirectional model becomes available. Training-and-validation wall-clock times additionally depend on the recipes and stopping rules in Table S2.

Statistical implementation

Bootstrap samples draw subjects with replacement and use the same sampled subjects for both configurations, retaining all nights of each sampled subject. The 95% percentile interval uses the 2.5th and 97.5th percentiles of 10,000 paired resamples. The primary Sleep-EDF analysis uses bootstrap seed 2026; context analyses use 2028, and the SHHS1 E1/E2 and E3–E2 extensions use 2031 and 2032, respectively. Each analysis retains its archived seed and comparison family. Two-sided Wilcoxon tests use `zero_method="wilcox"`, excluding zero differences from ranking, no continuity correction and automatic selection of the calculation method. If all paired differences are zero, the implementation returns $p = 1$. Holm correction applies to Wilcoxon p -values within each stated family, not to the percentile intervals. Non-significance is not interpreted as equivalence between configurations.

Supplementary results

Table S4: Per-class Sleep-EDF metrics for E3 and E6 (seed 42). Differences are calculated before rounding.

Stage	E3 Prec.	E3 Rec.	E3 F1	E6 Prec.	E6 Rec.	E6 F1	$\Delta F1$
W	.9338	.9271	.9304	.9391	.9169	.9279	+0.0026
N1	.5149	.5410	.5276	.5032	.5219	.5124	+0.0152
N2	.8583	.8447	.8514	.8401	.8368	.8385	+0.0130
N3	.8072	.8097	.8085	.7160	.7779	.7457	+0.0628
REM	.8274	.8413	.8343	.8227	.8197	.8212	+0.0131

Table S5: Pooled E3 confusion matrix on Sleep-EDF (rows: reference; columns: prediction).

	W	N1	N2	N3	REM
W	61,133	3,888	416	56	448
N1	3,283	11,643	4,867	105	1,624
N2	539	5,496	58,394	2,271	2,432
N3	30	45	2,377	10,558	29
REM	483	1,542	1,984	90	21,736

Table S6: Sleep-EDF pooled macro-F1 under two fixed seeds. Seed 42 is primary; seed 123 is a post-protocol sensitivity repeat. Differences are calculated before rounding.

Configuration	Seed 42	Seed 123	Difference
E0	.7754	.7755	+0.0000
E1	.7802	.7776	-.0026
E2	.7835	.7828	-.0007
E3	.7904	.7883	-.0022
E4	.7891	.7887	-.0004
E6	.7691	.7780	+0.0089

Table S7: Seed sensitivity of the protocol-specified Sleep-EDF contrasts. Bootstrap intervals describe differences in pooled macro-F1; Wilcoxon tests use paired subject-wise macro-F1. Holm correction is applied to the four tests within each seed, not to the marginal bootstrap intervals.

Contrast	Seed 42: Δ [95% CI], p_{Holm}	Seed 123: Δ [95% CI], p_{Holm}
E1–E0	.004811 [–.000915, .010193], .1022	.002188 [–.002870, .007043], .9130
E2–E1	.003251 [–.002370, .008808], .1036	.005143 [–.000571, .011292], .2400
E3–E2	.006962 [.000305, .014520], .8989	.005478 [–.002611, .015796], .9130
E3–E6	.021319 [.012179, .030698], .0012	.010249 [.002435, .017989], .1313

Table S8: Seed-123 wall-clock training and validation time (10 folds, sequential, same GPU). E1 time covers sequence training only because its encoder and cached features are reused from E0.

Configuration	Total	Mean per fold
E0: 15-CNN + BiLSTM	33 h 35 min 36 s	3 h 21 min 34 s
E1: 15-CNN + TCN	14 min 17 s	1 min 26 s
E2: ResNet-1D + TCN	2 h 44 min 57 s	16 min 30 s
E3: fixed scale	3 h 16 min 40 s	19 min 40 s
E4: band-pass	2 h 55 min 56 s	17 min 36 s
E6: record-wise z -score	2 h 24 min 44 s	14 min 28 s

Table S9: Seed-123 sensitivity and E4 extension on the same 180 SHHS1 participants. These are post-protocol results on the previously evaluated cohort.

ID	Subject-mean macro-F1	Pooled macro-F1	Accuracy	κ
E0	.5303	.5785	.6738	.5420
E2	.5409	.5858	.6874	.5597
E3	.5634	.6031	.7003	.5772
E4	.5732	.6147	.7064	.5869
E6	.5503	.5865	.6873	.5552

Table S10: Paired SHHS1 seed-123 extension. Differences and subject-cluster percentile bootstrap intervals concern subject-mean macro-F1 (10,000 resamples). Two-sided Wilcoxon p -values are Holm-adjusted within this four-contrast extension, separately from the primary seed-42 family. W/T/L denotes subjects with a positive/zero/negative paired difference.

Contrast	Δ subject-mean macro-F1	95% CI	p_{Holm}	W/T/L
E4–E2	+.0323	[.0236, .0410]	7.40×10^{-13}	135/1/44
E3–E4	–.0098	[–.0125, –.0073]	1.89×10^{-10}	56/1/123
E3–E6	+.0131	[.0023, .0239]	.0085	111/0/69
E3–E0	+.0331	[.0225, .0438]	5.19×10^{-9}	131/0/49

Table S11: Additional E4–E3 comparisons on the same 180 SHHS1 participants and 169,012 labelled epochs. Every result uses the mean probabilities of all ten source-fold models at its stated seed. Intervals are paired subject-cluster percentile bootstrap intervals (10,000 resamples, bootstrap seed 2031), not intervals over training seeds. E4 seed 42 was newly inferred from historical checkpoints; its E3 control and both seed-123 results are verified archived predictions. The original seed-123 Holm analysis remains in Table S10; no new seed-42 p -value is added to the primary comparison family.

Seed	Model	Subject-mean macro-F1	Pooled macro-F1	N3 recall	N3 F1
42	E3	.5680	.6099	.2582	.4055
42	E4	.5811	.6234	.2728	.4233
123	E3	.5634	.6031	.2410	.3840
123	E4	.5732	.6147	.2591	.4068
E4–E3, seed 42		+.0131	95% CI [.0111, .0153] for subject-mean difference		
E4–E3, seed 123		+.0098	95% CI [.0073, .0125] for subject-mean difference		

Table S12: Secondary and post-hoc SHHS1 paired contrasts (seed 42; 180 subjects). Intervals are marginal 95% paired subject-cluster percentile bootstrap intervals for the subject-mean difference, with 10,000 resamples. The E1/E2 extension was specified before their SHHS1 inference but after the cohort had been evaluated with E0/E3/E6. E3–E2 was specified after their separate marginal results had been examined.

Contrast	Δ subject-mean macro-F1	95% CI	p	W/T/L
E1–E0	+0.00651	[.00192, .01087]	.00325 ^a	108/0/72
E2–E1	-.01280	[–.02209, –.00350]	.01005 ^a	80/0/100
E3–E2	+0.04750	[.03724, .05792]	2.35×10^{-17} ^b	147/0/33

^aTwo-sided subject-wise Wilcoxon, Holm-adjusted within the two-contrast E1/E2 extension (bootstrap seed 2031). ^bUnadjusted two-sided Wilcoxon for one separately specified post-hoc comparison (bootstrap seed 2032); it is not part of either historical Holm family. The contrasts compare specified pipeline configurations, not isolated architectural effects.

Table S13: Per-class transfer of E3; Pred/True is the predicted-to-reference class count ratio, not a measure of probability calibration. Differences are calculated before rounding.

Stage	Sleep-EDF		SHHS1				
	F1	Pred/True	Prec.	Rec.	F1	Pred/True	Δ F1
W	.9304	.993	.8823	.8014	.8399	.908	-.0905
N1	.5276	1.051	.2528	.5437	.3451	2.151	-.1824
N2	.8514	.984	.7394	.7349	.7372	.994	-.1143
N3	.8085	1.003	.9440	.2582	.4055	.273	-.4030
REM	.8343	1.017	.6052	.8939	.7218	1.477	-.1125

Table S14: Pooled E3 confusion matrix on SHHS1 (rows: reference; columns: prediction).

	W	N1	N2	N3	REM
W	29,638	4,481	1,171	21	1,672
N1	874	3,807	601	3	1,717
N2	2,348	5,947	56,063	324	11,605
N3	78	39	16,674	5,888	127
REM	655	784	1,311	1	23,183

Table S15: Per-class SHHS1 metrics for E3 and E2 (seed 42). Differences are calculated before rounding.

Stage	E3 Prec.	E3 Rec.	E3 F1	E2 Prec.	E2 Rec.	E2 F1	Δ F1
W	.8823	.8014	.8399	.8780	.7275	.7957	+0.0442
N1	.2528	.5437	.3451	.2367	.3528	.2833	+0.0618
N2	.7394	.7349	.7372	.7362	.6937	.7143	+0.0228
N3	.9440	.2582	.4055	.9659	.2496	.3967	+0.0087
REM	.6052	.8939	.7218	.4886	.9450	.6442	+0.0776

Table S16: Reference-label oracle on E3 predictions (seed 42). Each row corrects only the indicated confusion-matrix cell. The percentage is relative to the observed pooled-score difference between Sleep-EDF out-of-fold evaluation and SHHS1 ensemble evaluation, not an isolated domain-shift effect or an estimate of achievable model performance.

Condition	macro-F1	Numerical gap recovered
Observed E3 on SHHS1	.6099	—
Correct N3→N2 only	.7444	.1345 (74.5%)
Correct N2→REM only	.6596	.0497 (27.5%)
E3 on Sleep-EDF	.7904	—

Table S17: Post-hoc reference-label boundary and stable-interior diagnostics. Seed 123 reuses the same subjects and label-defined regions; it is a sensitivity repeat, not an independent cohort. All metrics are pooled within each region.

Cohort	Seed	ID	Boundary macro-F1	Stable macro-F1	Boundary N3 recall	Stable N3 recall
Sleep-EDF	42	E0	.6179	.8175	.6121	.9216
Sleep-EDF	42	E3	.6321	.8346	.6326	.9316
Sleep-EDF	123	E0	.6181	.8173	.6174	.9162
Sleep-EDF	123	E3	.6343	.8299	.6385	.9285
SHHS1	42	E0	.4323	.5959	.0633	.4008
SHHS1	42	E3	.4689	.6220	.0733	.3900
SHHS1	123	E0	.4336	.5957	.0639	.3885
SHHS1	123	E3	.4605	.6166	.0656	.3660

For a label change between consecutive valid epochs $j - 1$ and j , where j is the first epoch with the new label, the boundary neighbourhood contains $\{j - 1, j, j + 1\}$ within the same contiguous recording segment. Overlapping neighbourhoods are combined. Unknown-label gaps and recording boundaries break contiguity. Stable-interior epochs are at least three epoch positions from *both* sides of every label change, using distance $\min(|i - (j - 1)|, |i - j|)$; segments without a change are stable. The two regions do not exhaust the scored epochs: Sleep-EDF contains 43,237 boundary, 130,693 stable and 21,539 remaining epochs; SHHS1 contains 44,744, 103,195 and 21,073, respectively. The remaining epochs are not simply the distance-two bin, because the boundary neighbourhood is centred on the first new-label epoch. These masks include all stage changes, not just N2–N3. Where persistent boundaries are reported in the archived diagnostics, each side must contain at least three consecutive epochs of its respective label. Region metrics retain the fixed five-class convention used in the main analysis.

E6 N3 recall is 0.2005 overall and 0.0721 in the same SHHS1 boundary neighbourhood (seed 42), compared with 0.2582 and 0.0733 for E3. These regions use reference labels on both sides and are offline diagnostics, not direct measurements of physiological transitions or deployable detectors. Differences in class composition between regions also preclude interpreting their score difference as an isolated effect of proximity to a label change. N3 under-detection persists in SHHS1 stable interiors in both seeds.

Supplementary analyses: representation and context

The representation analysis uses the same 1,000 held-out epochs per fold for E1 and E2: 200 per stage, sampled without replacement with a round-robin allocation across available subjects (sampling seed 42 + fold index). In each representation and fold separately, features are standardised and projected onto 20 principal components; these transforms are fitted to the sampled test features solely for this descriptive analysis, not fed back into the staging model. Silhouette uses Euclidean distance in the resulting space and the reference stage labels. The mean coefficient across ten folds is 0.2126 (SD 0.0356) for E1 and 0.1048 (SD 0.0229) for E2. E1 consists of 75 softmax-output features, whereas E2 is an unconstrained 128-dimensional representation; this difference and the analysis transforms prevent treating silhouette as a like-for-like measure of downstream predictive utility.

For the context ablation, the pretrained E0 encoder and 75-dimensional input interface are retained. Each removed feature dimension is replaced by its mean over valid labelled training epochs in that fold; the same replacement vector is applied to training, validation and test. TCN models are retrained for C, CP and CN with the specified E1 recipe and seed 42, while full CPN reuses E1. Thus the intervention changes the information available during both fitting and inference; it is not an inference-only deletion. The local protocol specifies the primary CPN–C boundary contrast and two secondary contrasts, with Holm correction across those three subject-wise Wilcoxon tests. Bootstrap intervals describe paired differences in pooled boundary macro-F1 under 10,000 subject-cluster resamples (seed 2028).

Table S18: Context-group ablation within E1 (seed 42). C, P and N denote current, previous and next epoch feature groups. Differences are calculated before rounding.

Condition	Overall macro-F1	Δ vs C	Boundary macro-F1	Δ vs C	N1 F1
C	.777477	—	.619102	—	.5011
CP	.777733	+0.000257	.617686	-.001416	.5022
CN	.781585	+0.004108	.618895	-.000207	.5055
CPN	.780230	+0.002753	.620055	+0.000953	.5115

For the protocol-specified CPN–C boundary comparison, $\Delta_{\text{macro-F1}} = 0.000953$, 95% CI $[-0.004588, 0.006568]$ and Holm $p = 1.000$ (36/78 subjects improved). Secondary CPN–CP and CPN–CN intervals also include zero with Holm $p = 1.000$. These comparisons did not detect an incremental boundary-performance benefit from the explicit P/N feature groups under the E1 recipe. The TCN continues to use temporal context across the sequence in every condition.




Figure S1: Supporting representation and context analyses. (a) Silhouette coefficients by test fold. (b) Paired 95% confidence intervals for pooled boundary macro-F1 differences: CPN minus C, CP and CN (Sleep-EDF, E1, seed 42). All intervals include zero; Holm-adjusted Wilcoxon $p = 1.000$ for each contrast. Plots were drawn from archived aggregates using repository code with OpenAI Codex assistance.

Supplementary Methods S2: additional intervention experiments

These experiments address errors observed in the original SHHS1 evaluation. They use the same 180 test participants. Model selection uses source validation, and adaptation uses unlabelled target data. All target recordings are processed in full, without annotation-based windowing, before applying the historical scoring mask (169,012 labelled epochs from 183,528 inferred epochs). Every contrast uses a control with the same inference context. In particular, the unmodified full-record E3 control is not interchangeable with the historical trimmed-context E3 ensemble. We report 10,000 paired participant-bootstrap resamples (seed 2031), with marginal 95% intervals, separately from the primary Holm families. These intervals condition on the fitted models and do not incorporate training variability.

Temperature scaling and prior correction: ten-fold ensemble

For each seed-123 E3 fold, a positive temperature is fitted by minimising negative log-likelihood on that fold's archived source-validation predictions, with bounds $[0.05, 20]$. This validation set was also used for historical checkpoint selection. The source prior is the mean calibrated validation posterior. EM estimates a target prior from 4,989 full-record epochs of five fixed, unlabelled adaptation participants, disjoint from the 180 test participants. Initialisation uses the source prior; the stopping tolerance is 10^{-8} , the maximum is 1,000 iterations, and a 10^{-12} floor prevents zero probabilities. No true target prevalence or class-specific multiplier is supplied. For each fold, the calibrated posterior is multiplied by the estimated target/source prior ratio and renormalised. All ten fits are frozen before target inference. We average all ten fold probabilities separately for the raw, temperature-only and temperature-plus-EM arms, without fold selection. Temperature scaling preserves each model's argmax but can alter the argmax after probability averaging across differently calibrated models.

Table S19: Additional calibration experiment: all ten E3 folds, seed 123, full-record inference on 180 SHHS1 participants. Each column averages the same ten models. A dash denotes undefined N3 precision because the EM ensemble predicts no N3 epochs.

Metric	Raw	Temperature	Temperature + EM
Subject-mean macro-F1	.5590	.5595	.4899
Pooled macro-F1	.5989	.5994	.5021
Accuracy	.6957	.6957	.6372
N3 precision	.9443	.9446	—
N3 recall	.2410	.2421	.0000
N3 F1	.3840	.3855	.0000
N2 recall	.7306	.7287	.7288

Temperature-only minus raw subject-mean macro-F1 is 0.00049 (95% CI $[0.00002, 0.00097]$), a small absolute change. Temperature-plus-EM minus temperature-only is -0.0696 ($[-0.0793, -0.0600]$). All ten EM fits converge, but each places its N3 prior near the numerical floor; 98.83% of reference N3 epochs are then labelled N2. The estimated prior describes the

output of this EM procedure on five adaptation participants, rather than their measured N3 prevalence. In this experiment, EM correction eliminated N3 predictions and reduced subject-mean macro-F1 relative to temperature scaling alone.

Class-weighted TCN: matched ten-fold ensemble

For each seed-123 fold, the corresponding verified historical E3 encoder is frozen and both sequence classifiers are trained with the same initial weights and source batch order. The control uses unweighted cross-entropy; the intervention uses $N/(5n_c)$ computed separately from that fold’s valid source-training labels. Both arms use Adam at learning rate 0.0005, batches of eight recordings, gradient norm clipping at 1, at most 300 epochs and patience 30. Strict maximum pooled source-validation macro-F1 selects each checkpoint. No target labels are used for fitting, stopping, class weights or selection. All 20 checkpoints are frozen before the new target ensemble evaluation, and all ten probabilities are averaged in each arm. The historical E3 model is not the matched unweighted control.

Verified CPU pairs from folds 0–5 are reused; folds 6–9 are fresh CUDA pairs, not continuations of partially trained CPU models. Both arms within each fold use the same backend. PyTorch versions are 2.5.1 (CPU) and 2.11.0 (CUDA); the experiment is therefore mixed-backend, not a uniform-backend replication. Independent CPU replay of CUDA source-validation and outer-test predictions reproduces their confusion matrices and metrics. New training timings do not replace the historical operational benchmark. This intervention trains sequence classifiers on frozen features, not the encoders end-to-end. The earlier fold-0 result is retained below, not counted as an independent replication of the ten-fold result.

Table S20: Matched class-weighting experiment: ten-fold ensembles, seed 123, full-record inference on 180 SHHS1 participants. Both arms average ten newly trained sequence classifiers with the same fold-specific frozen encoders. Per-class metrics and error rates are pooled.

Metric	Unweighted	Weighted
Subject-mean macro-F1	.5530	.5514
Pooled macro-F1	.5930	.5928
Accuracy	.6892	.6663
N3 precision	.9462	.9342
N3 recall	.2345	.3225
N3 F1	.3759	.4795
N2 recall	.7260	.6728
N3→N2 rate	.7497	.6549
N2→N3 rate	.0038	.0063

Weighted minus unweighted subject-mean macro-F1 is -0.0016 (95% CI $[-0.0049, 0.0017]$). For pooled metrics, the differences are -0.0002 ($[-0.0032, 0.0027]$) in macro-F1, 0.1036 ($[0.0938, 0.1137]$) in N3 F1, 0.0880 ($[0.0805, 0.0954]$) in N3 recall, -0.0120 ($[-0.0187, -0.0050]$) in N3 precision and -0.0531 ($[-0.0577, -0.0488]$) in N2 recall. N3→N2 errors decrease by 0.0948, while N2→N3 errors increase by 0.0025. Intervals that include zero do not establish equivalence or non-inferiority. Bootstrap resampling pairs participants, not epochs, and conditions on the fixed fitted models; it does not quantify training or backend variability.

On the source outer tests (78 participants, 195,469 labelled epochs, out-of-fold rather than ensembled predictions), unweighted/weighted pooled macro-F1 is 0.7878/0.7840 and accuracy 0.8330/0.8213. N3 recall increases from 0.8059 to 0.8724, whereas precision decreases from 0.8121 to 0.7404 and F1 from 0.8090 to 0.8010. Increased recall alone therefore does not establish an overall improvement on source or target.

Class-weighted TCN: earlier matched single-fold pilot

This comparison uses fold 0 and seed 123. The verified E3 encoder is frozen, and both TCNs are trained anew with identical initial weights and source batch order. The control uses unweighted cross-entropy; the intervention uses $N/(5n_c)$ from valid source-training labels only. Weights in W/N1/N2/N3/REM order are (0.5755, 1.8330, 0.5671, 3.2466, 1.5493). Both arms use Adam with learning rate 0.0005, batches of eight recordings, gradient norm clipping at 1, at most 300 epochs, and patience 30. Strict improvement in pooled source-validation macro-F1 determines selection. The unweighted/weighted runs stop at epochs 42/53 and select epochs 12/23. Target labels are not used for fitting, stopping, class-weight choice or checkpoint selection. The retrained unweighted TCN, not the historical E3 checkpoint, is the matched control.

Table S21: Matched class-weighting experiment: fold 0, seed 123, 180 SHHS1 participants. Both columns are single-model predictions, not ten-fold ensembles.

Metric	Unweighted	Weighted
Subject-mean macro-F1	.5431	.5363
Pooled macro-F1	.5861	.5790
Accuracy	.6837	.6510
N3 precision	.9400	.9253
N3 recall	.2432	.3221
N3 F1	.3864	.4779
N2 recall	.7646	.6708

Weighted minus unweighted subject-mean macro-F1 is -0.0068 (95% CI $[-0.0135, 0.0001]$), whereas N3 recall increases by 0.0789 ($[0.0720, 0.0863]$) and N2 recall decreases by 0.0938 (difference interval $[-0.1097, -0.0785]$). The first interval does not establish equivalence or non-inferiority. On the source outer-test set (eight participants, 19,506 labelled epochs), pooled macro-F1 is $0.7716/0.7672$ and N3 recall is $0.8678/0.9098$ for unweighted/weighted training. N3 F1 decreases from 0.8031 to 0.7698 on source, showing that greater recall alone is not sufficient evidence of a better classifier.

ADAST: limited-budget matched ten-fold adaptation comparison

Ten new CUDA pairs use the same harmonised ADAST design as the earlier pilot below, on all source outer folds at seed 123. The upstream model/configuration files and similarity penalty are unchanged. Both arms in every fold start from identical weights and receive the same source permutations and 1,140 updates. Each epoch samples 4,864 valid source-training epochs, not a complete pass through that fold’s source data. The same five unlabelled adaptation participants supply 4,989 epochs; the 180 test participants provide no training signals or labels. All 20 final fixed-budget checkpoints are frozen before local target inference or scoring. For each arm, maximum dual-head logits are softmaxed within each fold and all ten probability arrays are averaged. Target inference uses source attention for the source-only arm and target attention for ADAST. The historical full-record scoring mask is then applied to both arms. Source results use each outer fold’s test subjects once (78 participants, 195,469 valid epochs), with source attention for both arms; they are out-of-fold predictions, not ensembles. The CUDA pairs and CPU inference are kept separate from the historical operational benchmark. Participant-bootstrap intervals condition on these fitted models, not training randomness.

Table S22: Matched ten-fold ADAST comparison, seed 123: 180 SHHS1 participants, 169,012 valid scored epochs. Both arms use ten fresh CUDA-trained models and local full-record inference; the source-only arm is the matched dual-head control, not E3.

Metric	Source-only	ADAST
Subject-mean macro-F1	.4949	.4497
Pooled macro-F1	.5539	.5072
Accuracy	.6977	.6951
Cohen’s κ	.5488	.5367
N3 precision	.9266	.9225
N3 recall	.3195	.4269
N3 F1	.4752	.5837
N1 recall	.1974	.0000
N2 recall	.8712	.8655
REM recall	.4343	.2938
N3→N2 rate	.6747	.5588

ADAST minus source-only subject-mean macro-F1 is -0.0452 (95% CI $[-0.0558, -0.0347]$), whereas the N3-recall difference is 0.1074 ($[0.0714, 0.1422]$) and the N3-F1 difference is 0.1085 ($[0.0721, 0.1442]$). N3→N2 errors decrease, but more than half of true N3 epochs remain misclassified as N2. The ADAST ensemble produces no N1 predictions, and REM recall decreases. On source out-of-fold testing, source-only/ADAST pooled macro-F1 is $0.6665/0.5276$, accuracy is $0.7724/0.7116$, N3 recall is $0.7689/0.5933$, N1 recall is $0.1538/0.0039$ and REM recall is $0.5895/0.2337$. The source subject-mean macro-F1 difference is -0.1388 ($[-0.1520, -0.1250]$). Thus, target N3 recall improved while aggregate source and target performance declined. The comparison uses a shared architecture and training budget for ADAST and its source-only control. The subsequent source-validation studies examine training-budget and loss-component effects; the full-source ten-fold evaluation below tests transfer under the larger budget.

ADAST: earlier matched single-fold pilot

The harmonised implementation uses the upstream ADAST model classes and similarity penalty at revision [e0fb503](#), with this study’s E3 preprocessing and source split (fold 0, seed 123). It is not an exact reproduction of the published data protocol. Both arms use the same initialisation, source permutations and 1,140 updates: two rounds of 15 epochs, each with 38 batches of 128 epochs. The shorter unlabelled-target loader determines this budget; each epoch presents 4,864 of the 157,200 source-training epochs, not a full source pass. The five adaptation participants supply 4,989 epochs. The matched source-only control retains the two classifiers and similarity regularisation, but has no target forward pass, adversarial loss or pseudo-label loss. Both arms use the same source-loss schedule (1 then 0.1), Adam (learning rate 0.001, betas 0.5/0.99, weight decay 0.0003). After epoch 10 of round 0, the learning rate is multiplied by 0.1 for every joint-optimizer parameter group in the implemented adapter, not only the encoder. ADAST uses adversarial weight 1, similarity weight 0.001, and target-loss weights 0 then 0.01. Pseudo-labels are fixed at each round’s start from the mean of the two logits. Inference takes the elementwise maximum of those logits, with source attention for the control and target attention for ADAST on SHHS. The final fixed-budget checkpoints are used; no target-label validation or performance-based selection is performed.

Table S23: Harmonised ADAST comparison: fold 0, seed 123, 180 SHHS1 participants. The source-only arm is the matched dual-head control, not E3.

Metric	Source-only	ADAST
Subject-mean macro-F1	.4755	.4214
Pooled macro-F1	.5319	.4766
Accuracy	.6820	.6762
N3 precision	.9277	.9270
N3 recall	.2657	.2622
N3 F1	.4131	.4087
N1 recall	.2521	.0077

ADAST minus source-only subject-mean macro-F1 is -0.0542 (95% CI $[-0.0641, -0.0443]$); the N3-recall difference is -0.0035 ($[-0.0426, 0.0345]$). On the source outer-test set, evaluated with source attention in both arms, pooled macro-F1 is 0.6562/0.5015 and N1 recall is 0.1049/0.0015. Thus the deterioration is not restricted to SHHS. This single-fold, limited-budget comparison does not rank ADAST against the ten-fold E3 ensemble or establish that unsupervised adaptation is generally ineffective. This pilot is retained separately from the new ten-fold CUDA comparison. Subsequent source-validation budget and loss-component experiments are reported below.

ADAST: source-validation learning-budget and loss-component experiments

Two subsequent development campaigns used source fold 0, seed 123: source-only/ADAST at limited/full-source learning budgets, then a fresh full-source ADAST reference with three single-factor loss interventions. Training and source-validation data comprise 157,200 and 18,763 valid epochs respectively. Both campaigns reuse the same five unlabelled adaptation participants (4,989 epochs), without source outer-test or SHHS test access. Each model trains for two rounds of 15 epochs on T4 CUDA float32, using matched initialisation, source order and the earlier optimiser schedule. Limited training has 38 updates per epoch (1,140 total); full-source training has 1,229 updates per epoch (36,870 total), retaining the final batch of 16 source epochs. The larger budget also increases target presentations and changes scheduler timing in update units; it is not a source-coverage-only intervention.

Both attention paths are evaluated on the same source validation set. The selected checkpoint maximises source-attention validation macro-F1, retaining the first tie; epoch-30 checkpoints are reported separately. No target-attention validation rule or SHHS score selects a checkpoint. Full-validation CPU inference from all eight selected/final checkpoints in each campaign and recomputation of all 124 initial/per-epoch validation diagnostics are performed independently. Participant-bootstrap intervals from the ten-fold transfer experiments are not attached to these one-fold development comparisons.

Table S24: Source-validation learning-budget comparison: fold 0, seed 123. All scores use source attention; selected checkpoints follow the fixed validation criterion.

Arm	Budget	Selected epoch	Selected macro-F1	Selected N1 recall	Final macro-F1	Final N1 recall
Source-only	Limited	15	.6585	.1956	.6316	.1243
ADAST	Limited	11	.6470	.1673	.5236	.0000
Source-only	Full source	9	.7174	.5380	.6911	.2847
ADAST	Full source	22	.6934	.2338	.6784	.1591

The larger budget improves selected-checkpoint macro-F1 in both learning conditions and avoids complete N1 loss at epoch 30 in ADAST. Nevertheless, selected-checkpoint source-only macro-F1 remains higher. For the full-source pair at epoch

30, ADASt/source-only N3 recall is .8068/.5783, whereas N1 recall is .1591/.2847. In the limited ADASt arm, N1 recall falls .1060 $\rightarrow$.0096 at epochs 15 $\rightarrow$ 16, when source CE weight falls 1 $\rightarrow$.1 and pseudo CE weight rises 0 $\rightarrow$.01 simultaneously. In the full-source ADASt reference, N1 recall instead rises .1660 $\rightarrow$.1934 at that boundary. The limited-budget trajectory is not a universal epoch-16 collapse.

The second campaign keeps the full-source budget and changes one configured component per arm: source CE remains 1 in round two, pseudo CE remains 0 in both rounds, or adversarial weight is 0. The latter arm skips discriminator forward/backward/optimiser steps, retaining target forward passes, the pseudo-label schedule and other losses. The discriminator weights remain unchanged and its optimiser state remains empty. The reference, keep-source and no-pseudo arms have identical saved validation logits through epoch 15; their interventions take effect from epoch 16. The fresh reference reproduces all 30 validation metrics of the previous full-source ADASt arm. Similarity weight remains .001 throughout.

Table S25: Full-source ADASt loss interventions on source validation: fold 0, seed 123. Selected and final checkpoints are separate comparisons. The selected no-pseudo checkpoint precedes intervention activation and is also present in the reference trajectory.

Arm	Checkpoint	Epoch	macro-F1	N1 recall	N3 recall	REM recall
Reference	Selected	22	.6934	.2338	.8261	.7139
Keep source CE	Selected	22	.6941	.2664	.8323	.6490
No pseudo CE	Selected	7	.6891	.4902	.8834	.5486
No adversarial loss	Selected	29	.6943	.2416	.7115	.6865
Reference	Final	30	.6784	.1591	.8068	.7544
Keep source CE	Final	30	.6809	.2064	.8095	.6742
No pseudo CE	Final	30	.6717	.1769	.8192	.6430
No adversarial loss	Final	30	.6777	.2425	.6025	.6764

Retaining source CE increases selected-checkpoint N1 F1 .2988 $\rightarrow$.3238 and recall .2338 $\rightarrow$.2664, while N1 precision is .4138/.4128. N3 recall increases slightly but N3 precision decreases .7533 $\rightarrow$.7399 and F1 .7880 $\rightarrow$.7834; REM recall and F1 decrease. Its selected macro-F1 gain over the reference is .000673. Removing adversarial loss gives the highest selected macro-F1 among these four arms, a gain of .000887, but lowers N3 recall by .1146 and raises N3 precision .7533 $\rightarrow$.8294. At fixed epoch 30, no-pseudo improves N1 recall modestly but lowers overall macro-F1 and REM recall; no-adversarial increases N1 recall but lowers N3 recall by .2043. High N1 recall at the selected no-pseudo epoch 7 is not an effect of disabling pseudo-label learning, which has zero weight in the reference at that point. These interventions identify distinct stage-wise trade-offs rather than a uniformly improved configuration.

Table S26: Five-stage F1 for the full-source ADASt loss interventions on source validation.

Arm	Checkpoint	W	N1	N2	N3	REM
Reference	Selected	.8771	.2988	.8151	.7880	.6880
Keep source CE	Selected	.8831	.3238	.8100	.7834	.6700
No pseudo CE	Selected	.8818	.4055	.7692	.7776	.6116
No adversarial loss	Selected	.8804	.3111	.8205	.7660	.6934
Reference	Final	.8812	.2295	.8073	.7944	.6796
Keep source CE	Final	.8835	.2722	.8014	.7880	.6596
No pseudo CE	Final	.8797	.2465	.7925	.7848	.6548
No adversarial loss	Final	.8772	.3074	.8108	.7173	.6757

Table S27: Target-attention diagnostics of the same checkpoints on the same source validation set, not SHHS evaluation. Target attention does not select checkpoints.

Arm	Checkpoint	macro-F1	N1 recall	N3 recall
Reference	Selected	.6780	.2225	.8302
Keep source CE	Selected	.6789	.2599	.8599
No pseudo CE	Selected	.6803	.4259	.8551
No adversarial loss	Selected	.5227	.1521	.0773
Reference	Final	.6674	.1586	.8116
Keep source CE	Final	.6628	.1899	.8682
No pseudo CE	Final	.6546	.1517	.8116
No adversarial loss	Final	.5026	.1352	.0628

Without adversarial loss, the target-attention path receives neither adversarial nor pseudo-CE loss contributions in round one; source CE uses source attention, and similarity regularises the two classifier heads. Target forward passes and optimiser/weight decay are retained. The no-adversarial path therefore has a different learning condition, as intended by the intervention. Its poor target-attention N3 recall in Table S27 must be considered alongside the stronger source-attention scores; it is not evidence of better cross-cohort adaptation.

The selected source-validation ranking is no-adversarial, keep-source, reference, then no-pseudo by selected validation macro-F1; all remain below the full-source source-only score .7174. These source-validation comparisons distinguish the effects of loss components from the ten-fold SHHS transfer results reported above.

ADAST: three-arm confirmation on another source fold

Fold 1, seed 123, compared fresh source-only, ADAST-reference and retained-source-CE models with identical source orders; the ADAST arms shared target orders and round-one predictions. Each arm used 30 full-source passes over 152,825 training epochs: 1,194 updates per pass (batch 128, final batch 121), or 35,820 updates. Validation comprised 23,881 epochs (W/N1/N2/N3/REM: 9,912/1,911/7,133/1,968/2,957). The same five unlabelled adaptation participants were reused, without source outer-test or SHHS test access. Selection retained the source-validation rule above; epoch 30 was reported separately.

The comparison reports aggregate macro-F1 and all five stage F1 scores for both attention paths, using source attention for checkpoint selection. Selected and final checkpoints are reported separately to distinguish the best validation result from the end-of-training result.

Table S28: Fold-1 confirmation, seed 123: selected and final source-attention metrics on the same source validation. Stage columns are F1.

Arm	Checkpoint	Epoch	macro-F1	W	N1	N2	N3	REM
Source-only	Selected	10	.7370	.9432	.3200	.8427	.8568	.7224
Reference	Selected	27	.6907	.9392	.2175	.8021	.8096	.6852
Keep source CE	Selected	27	.6892	.9382	.2302	.7935	.8161	.6680
Source-only	Final	30	.7176	.9444	.2415	.8333	.8503	.7184
Reference	Final	30	.6857	.9391	.2175	.7968	.8077	.6672
Keep source CE	Final	30	.6814	.9397	.2158	.7865	.7999	.6653

Table S29: Target-attention diagnostics for the fold-1 checkpoints on the same source validation, not SHHS. This path does not select checkpoints. Stage columns are F1.

Arm	Checkpoint	macro-F1	W	N1	N2	N3	REM
Source-only	Selected	.6914	.8955	.2911	.8201	.7880	.6624
Reference	Selected	.6817	.9339	.2089	.7937	.8003	.6715
Keep source CE	Selected	.6727	.9318	.2085	.7802	.7864	.6564
Source-only	Final	.6213	.8318	.2052	.7650	.6403	.6642
Reference	Final	.6824	.9360	.2115	.7949	.7951	.6743
Keep source CE	Final	.6714	.9343	.2073	.7761	.7731	.6660

Retaining source loss increased selected N1 recall .1502 $\rightarrow$.1627, but reduced REM recall .7376 $\rightarrow$.6980. N3 recall changed .9644 $\rightarrow$.9649 and precision .6975 $\rightarrow$.7070. Through target attention, N3 recall increased .9670 $\rightarrow$.9756 but precision decreased .6826 $\rightarrow$.6587, lowering F1. At epoch 30, N1 F1 and macro-F1 were also lower than the reference: the selected N1 benefit was not uniform across checkpoints or paths. The fold-0 N1 gain and REM trade-off remain separate. Training/validation used 2,975.78 seconds on T4, excluding upload and result verification.

ADAST: full-source ten-fold transfer and matched budget contrast

The full-source comparison comprises ten source-only/ADAST pairs at seed 123. Four previously evaluated models from folds 0–1 are reused after checking their checkpoint, input and verification hashes; the retained-source-CE and other ablation arms are excluded. Sixteen models for folds 2–9 complete 30 epochs on T4 CUDA. Each epoch covers the entire source-training partition for that fold, yielding 35,670–37,620 updates per model over 30 epochs. The paired arms share initial weights and source permutations and use the same five unlabelled adaptation participants (4,989 epochs). The input bundle physically contains the source records covered by the earlier authorised fold-0/1 train/validation union; each model uses only its own fold’s training and validation roles. SHHS evaluation signals and labels remain local and are not used for training.

The main comparison uses epoch 30 in both arms and both budgets, rather than choosing a checkpoint on SHHS. Within each fold, the maximum of the two classifier logits is softmaxed; probabilities are averaged in float64 across ten folds before

float32 storage. SHHS uses source attention for source-only and target attention for ADAST. The full recordings supply 183,528 input epochs and the common scoring mask retains 169,012 epochs from 180 participants. Sleep-EDF outer-test predictions use source attention for both arms, each participant in its own fold; these are out-of-fold results, not ensemble predictions.

The new training outputs pass independent recomputation of all 496 validation diagnostics and CPU inference of all 32 selected/final checkpoints on their complete source-validation partitions through both attention paths. The local evaluation recomputes prediction hashes, confusion matrices and participant bootstrap results, with 40 target model-record and 20 source model-batch inference replays. These checks supplement the proof checks for the four reused models.

Table S30: Epoch-30 full-source pooled class metrics. SHHS1: 180 participants, 169,012 scored epochs, ten-model ensemble per arm. Sleep-EDF: 78 participants, 195,469 outer-test epochs, out-of-fold predictions. SO is the matched source-only dual-head control; AD is ADAST. Supports are true-class epoch counts.

Cohort	Arm	Stage	Precision	Recall	F1	Support
SHHS1	SO	W	.8317	.8372	.8344	36,983
	SO	N1	.1267	.1568	.1401	7,002
	SO	N2	.7152	.7981	.7544	76,287
	SO	N3	.9463	.2432	.3869	22,806
	SO	REM	.5631	.6976	.6231	25,934
	AD	W	.8216	.8441	.8327	36,983
	AD	N1	.1299	.1057	.1165	7,002
	AD	N2	.8103	.6369	.7132	76,287
	AD	N3	.8823	.6111	.7221	22,806
	AD	REM	.4375	.8360	.5744	25,934
Sleep-EDF	SO	W	.9061	.9293	.9176	65,941
	SO	N1	.4392	.2850	.3457	21,522
	SO	N2	.7892	.8878	.8356	69,132
	SO	N3	.8209	.7868	.8035	13,039
	SO	REM	.7362	.6727	.7030	25,835
	AD	W	.9216	.9034	.9124	65,941
	AD	N1	.4374	.2255	.2976	21,522
	AD	N2	.7580	.8429	.7982	69,132
	AD	N3	.6864	.8995	.7786	13,039
	AD	REM	.6295	.6278	.6287	25,835

Table S31: Matched participant-bootstrap budget contrasts on the same 180 SHHS1 participants. Effects are ADAST minus its source-only control; the difference of effects is full-source minus limited-budget. All four systems use the same 10,000 participant resamples (seed 2031). Intervals condition on the fitted models.

Contrast / metric	Difference	95% CI
Limited: subject-mean macro-F1	-.0452	[-.0558, -.0347]
Full-source: subject-mean macro-F1	.0445	[.0352, .0539]
Difference of effects: subject-mean macro-F1	.0897	[.0767, .1023]
Full-source: pooled macro-F1	.0440	[.0345, .0536]
Full-source: accuracy	-.0025	[-.0140, .0088]
Full-source: N1 F1	-.0236	[-.0385, -.0077]
Full-source: N2 F1	-.0412	[-.0549, -.0281]
Full-source: N3 F1	.3351	[.3025, .3683]
Full-source: REM F1	-.0488	[-.0629, -.0339]
Full-source: N3 recall	.3679	[.3442, .3912]
Full-source: N3→N2 rate	-.3836	[-.4076, -.3590]
Full-source: N2→N3 rate	.0192	[.0146, .0245]

Full-source subject-mean macro-F1 is .5010/.5455 and pooled macro-F1 .5478/.5918 for source-only/ADAST. N3→N2 errors fall from 17,163 to 8,414 of 22,806 true N3 epochs (rates .7526/.3689), while N2→N3 errors rise from 299 to 1,763 of 76,287 true N2 epochs (.0039/.0231). The N3 benefit is accompanied by lower N1, N2 and REM F1. REM recall increases, but its precision decreases enough to lower F1; recall alone therefore does not demonstrate a general REM improvement. Accuracy changes .6898 → .6873, with a paired interval spanning zero, not evidence of equivalence.

On Sleep-EDF, source-only/ADAST subject-mean macro-F1 is .6613/.6281 (difference -.0332, 95% CI [-.0434, -.0229]), pooled macro-F1 .7211/.6831 and accuracy .8003/.7707. N3 recall increases while its precision and F1 decrease. Thus the positive SHHS class-balanced effect does not represent uniform improvement across classes or cohorts. Increasing the budget also increases target exposure and changes the scheduler's timing in update units; the contrast does not isolate source-data

coverage. Folds and epochs are not treated as independent statistical samples, and the intervals do not include seed-to-seed training variation. The historical latency and workflow benchmarks are unchanged by this CUDA campaign.

ADAST: secondary source-validation checkpoint evaluation

Each of the same twenty full-source models supplies its checkpoint with the highest source-validation macro-F1 through source attention, using the first epoch in a tie. The SHHS attention paths, ten-fold probability averaging and scoring mask are unchanged; source results remain out-of-fold through source attention. No model is retrained or selected by SHHS performance. This secondary analysis changes the checkpoint rule, not the participants, adaptation data, loss or training budget. Epoch 30 remains the primary comparison across the two learning budgets.

Table S32: Secondary full-source evaluation with checkpoints selected by source validation. SHHS1 uses ten-fold ensembles on the same 180 participants; Sleep-EDF uses OOF predictions for 78 participants. Differences are ADAST minus source-only. The 10,000 participant-bootstrap draws (seed 2031) condition on fitted models.

Cohort	Arm	Subject-mean macro-F1	Pooled macro-F1	Accuracy
SHHS1	Source-only	.5073	.5552	.6818
	ADAST	.5396	.5854	.6709
Sleep-EDF	Source-only	.6681	.7237	.7947
	ADAST	.6368	.6921	.7730

Cohort	Contrast metric	Difference	95% CI
SHHS1	Subject-mean macro-F1	.0322	[.0234, .0412]
	Pooled macro-F1	.0302	[.0212, .0395]
	N3 recall	.3267	[.3050, .3483]
	N3 F1	.2892	[.2590, .3200]
Sleep-EDF	Subject-mean macro-F1	−.0313	[−.0416, −.0205]
	Pooled macro-F1	−.0316	[−.0416, −.0221]

Table S33: Secondary full-source pooled class metrics at source-validation-selected checkpoints. SHHS1: 169,012 scored epochs; Sleep-EDF: 195,469 outer-test epochs. SO is source-only; AD is ADAST. Supports are true-class epoch counts.

Cohort	Arm	Stage	Precision	Recall	F1	Support
SHHS1	SO	W	.8367	.8253	.8310	36,983
	SO	N1	.1276	.1807	.1495	7,002
	SO	N2	.7278	.7555	.7414	76,287
	SO	N3	.9333	.2828	.4340	22,806
	SO	REM	.5302	.7466	.6201	25,934
	AD	W	.8256	.8374	.8314	36,983
	AD	N1	.1271	.1231	.1251	7,002
	AD	N2	.8112	.5992	.6893	76,287
	AD	N3	.8891	.6095	.7232	22,806
	AD	REM	.4163	.8464	.5581	25,934
Sleep-EDF	SO	W	.9189	.9120	.9154	65,941
	SO	N1	.4197	.3546	.3844	21,522
	SO	N2	.8090	.8580	.8328	69,132
	SO	N3	.7565	.8182	.7861	13,039
	SO	REM	.7200	.6805	.6997	25,835
	AD	W	.9231	.9024	.9126	65,941
	AD	N1	.4229	.2605	.3224	21,522
	AD	N2	.7735	.8341	.8026	69,132
	AD	N3	.7062	.8875	.7865	13,039
	AD	REM	.6248	.6485	.6365	25,835

For source-only/ADAST on SHHS1, N3→N2 errors are 16,214/8,290 of 22,806 true N3 epochs (rates .7110/.3635); N2→N3 errors are 383/1,638 of 76,287 true N2 epochs (.0050/.0215). N1, N2 and REM F1 decrease despite the positive aggregate class-balanced effect. The source outer-test loss also persists. Both checkpoint rules therefore show an N3 benefit with stage/cohort costs. They reuse the same cohorts and related fitted models, not independent replication samples.

Provenance

The historical SHHS protocol matches the original run-manifest hash ('165d7cdf...fe93'), not the expanded post-run audit ('9541e233...fe9'). Table [S17](#) reuses seed-123 diagnostics matched to the EDF campaign and 20 E0/E3 fold-prediction files, and the SHHS1 extension manifest and 360 E0/E3 ensemble files. The retained files are linked to these analyses by their SHA-256 hashes. SHHS data and predictions remain externally held; E6 uses retained predictions. Snapshot reconciliation is documented in `Reports/SHHS_PROTOCOL_PROVENANCE.md`.